

F*ckUp-Free Artificial Intelligence: 100% Free, Offline, 8GB RAM — The Alternative That Doesn't Crash on Big Tasks

By Jose Solorzano Luna — Built on UFO v6.4 Cathedral Edition | Merges Windows Automation + Formal Safety | 5 Local Models, No API, No Credits

The Hook — What You Want to Hear First

- **100% Free** (USUALLY ■ EXPENSIVE) — No \$200/mo Operator / Claude Computer Use / Perplexity Computer / Grok / Kimi subscription. Unlimited use.
- **AI can see ■■■■■■ (your computer screen) + videos** — Learns YouTube tutorials, understands your screen via llava:7b Eyes.
- **AI can click or type for you (ask what you want)** — Windows automation, Excel, browser, any app.
- **(Available Offline) : (Works Offline) with 8GB RAM or more** — The worldwide-known way. Your screen never leaves your PC. 100% local via Ollama.

Why 'F*ckUp-Free'? Because other AIs f*ck up on big tasks and you pay for it. Mine is literally free OF fuckups. Play on words = more traffic, but technically true: fail-closed, hash-chained, replayable.

The One Sentence That Converts Curiosity to Need

Other AIs (Operator, Claude Computer Use, Perplexity Computer, Grok, Kimi) crash after a big task because their brain window is ~4k tokens. They forget step 1 by step 15, loop forever, and you can't prove what they did or view the underlying issues. Mine is worked around by auto-chunking, anti-loop Gate, and immutable Ledger with overview + underlying issue viewer.

The Problem With Other Artificial Intelligence Options

- **No formal safety.** One model decides and executes. If it loops, it loops forever. No Watcher-A/B, no Council, no Gate.
- **4k token window crash.** Giant tasks >500 chars or >10 steps overflow. They forget. No automatic splitting into fresh 4k windows, no files/phase1.txt memory carry.
- **No immutable history.** You can't prove what it did. Logs are mutable, editable, deletable.
- **No replay.** You can't verify a tutorial it learned. No Gnosis, no RCA.
- **Expensive + cloud.** \$20-\$200/mo, credits, API limits, screen sent to cloud. Not private, not offline.
- **No overview of issues.** No way to see what was processed, what was skipped, what underlying pronunciation errors caused missteps.

The Solution — Best of Both Worlds (UFO + Sentinel Substrate Merge)

I merged scrappy consumer UFO (works on 8GB) with formal military-grade Cathedral architecture. You get endless possibilities, all free, with safety others don't have.

Layer L0-L1: World Input / Model Ports

5 local Ollama models: phi3:mini (FAST - 1 click), llama3.2:3b (MID - 2-4 steps), llama3.1:8b (SMART - anti-loop/planning/tutorial), qwen2.5:7b (QWEN for Excel/code), llava:7b (EYES). All proposals only, no direct execution. Only AEGIS may mutate state.

Layer L2-L5: Epistemic Stratum — The Best of Both Worlds

- **VECTOR (L2):** Read-only sensing. Cannot mutate.
- **Watcher-A (L3):** Internal consistency auditor — detects loops, giant 4k token overflow, destructive commands (rm -rf, format). Cannot execute. Cannot bypass Gate.
- **Watcher-B (L4):** Adversarial auditor — audits Watcher-A itself for bias drift, audit corruption, silent failure, false confidence. This is key: most agents trust their own auditor. We don't.
- **Council Resolver (L5):** Merges A+B into exactly ONE SignalNode per cycle (Invariant I4). Agreement = reinforced confidence, Disagreement = escalated severity + reduced confidence, Critical disagreement = CRITICAL + minimum confidence. This is my TEAM VOTE, now formalized.
- **Signal Algebra:** Severity INFO/WARNING/ERROR/CRITICAL + confidence 0-1 + decay + escalation. Signals annotate, signals never decide (Signals annotate. Signals NEVER decide).

Layer L10-L12: Execution Stratum

- **CGIR (L10):** Causal Graph IR. Every task becomes StateNode (snapshot) + EventEdge (action + invariant mask + signal binding) + SignalNode. Not free text. Machine-checkable.
- **Gate (L11):** Pure deterministic function, only decision authority. Priority: 1 Ledger violation -> HARD_FAIL, 2 Critical Signal -> HARD_FAIL, 3 Temporal/hardware -> THROTTLE, 4 Contract violation -> HARD_FAIL, 5 VECTOR risk -> MODIFY/FALLBACK, 6 Planner quality -> ALLOW. Fail-closed, safe.
- **AEGIS/CESK (L12):** Only mutator. Execution = LOAD -> PROPOSE -> CHECK -> COMMIT -> PROVE. No other layer can mutate state. COMMIT only after Gate ALLOW/MODIFY. PROVE verifies state_after.

Layer L13-L14: Truth & Replay

- **Ledger/Chronicle (L13):** Every committed edge is hash-chained: $\text{hash}(n) = \text{SHA256}(\text{prev_hash} + \text{data})$. Immutable (I6 - no modification/deletion), tamper-evident. If you close tab, it's still on disk, provably. Solves losing content when closing tab.
- **Replay/Gnosis (L14):** Type 'replay' and it replays ledger, verifies chain (I9 Replayability). You can prove a YouTube tutorial was learned correctly. Root-cause analysis (RCA) + divergence detection.
- **Layer: Shadow/Ghost Mode:** For tutorial learning, run 2 shadows in parallel: 1s DETAILED (reliable) vs 3s FAST (unreliable). Neither mutates live state. Compare coherence, causal lift, risk, latency, cost, contract compliance, failure rate. Gate promotes best via Shadow -> Replay -> Value

Contract -> Gate Promotion. You don't choose 1s vs 3s, Gate does.

Step-by-Step Setup — So You Can Set It Up Yourself

1. **Install Ollama, pull models:** `ollama pull phi3:mini, llama3.2:3b, llama3.1:8b, qwen2.5:7b, llama:7b`
2. **Clone UFO:** `git clone https://github.com/microsoft/UFO ...` (UFO repo)
3. **Create venv ufo_env310 with requests, pillow:** `python -m venv ufo_env310; .\ufo_env310\Scripts\pip install requests pillow`
4. **Place ufo64.py (v6.4 Cathedral) in UFO folder:** Move downloaded file to `C:\Users\...\UFO\ufo64.py`
5. **Run:** `.\ufo_env310\Scripts\python.exe .\ufo64.py`
6. **Type help** -> shows all commands, so you never have to remember 'fastlearn tutorial named ...'
7. **Type status / replay / ledger** -> verifies immutable chain, 10-min auto-save reminder, safe Ctrl+C exit
8. **Learn:** learn tutorial named `my_first_test` -> 1s DETAILED default, or fastlearn tutorial named `X` -> 3s FAST

Why This Is By Far a Great Alternative — Endless Possibilities, All Free

- **Fixes 4k crash:** AI classifies if task is GIANT, splits into fresh 4k windows, carries memory via `files/phase1.txt`. No forgetting step 1 by step 15.
- **Fixes looping:** Watcher-A detects same task in last 5, Gate `HARD_FAIL` stops loop. Other AIs loop forever.
- **Overview + underlying issues:** status shows tutorials/screenshots/tasks/MB + ledger validity. replay shows every execution + hash. `pronunciation_errors_report.md` shows what was misheard + Team Vote consensus.
- **Affordable:** Works on 8GB RAM computer worldwide-known way. No \$3000 rig, no cloud.
- **Best of both worlds:** Scrappy automation + formal safety = capabilities not available elsewhere at any price.
- **Monetize:** Free on birthdays (Jan 29), paid rest via Quora upvotes/views. People would pay to learn how to set up what we got because it's so cool and affordable.